Introduction to Reinforcement Learning

Prof. Chia-Yu Lin

Yuan Ze University

2022 Spring

Thanks to the slides from Prof. Hung-Yi Lee in NTU.

Supervised Learning → RL

Human label

Cat

Human label

"3-3"?

It is challenging to label data in some tasks.

...... machine can know the results are good or not.
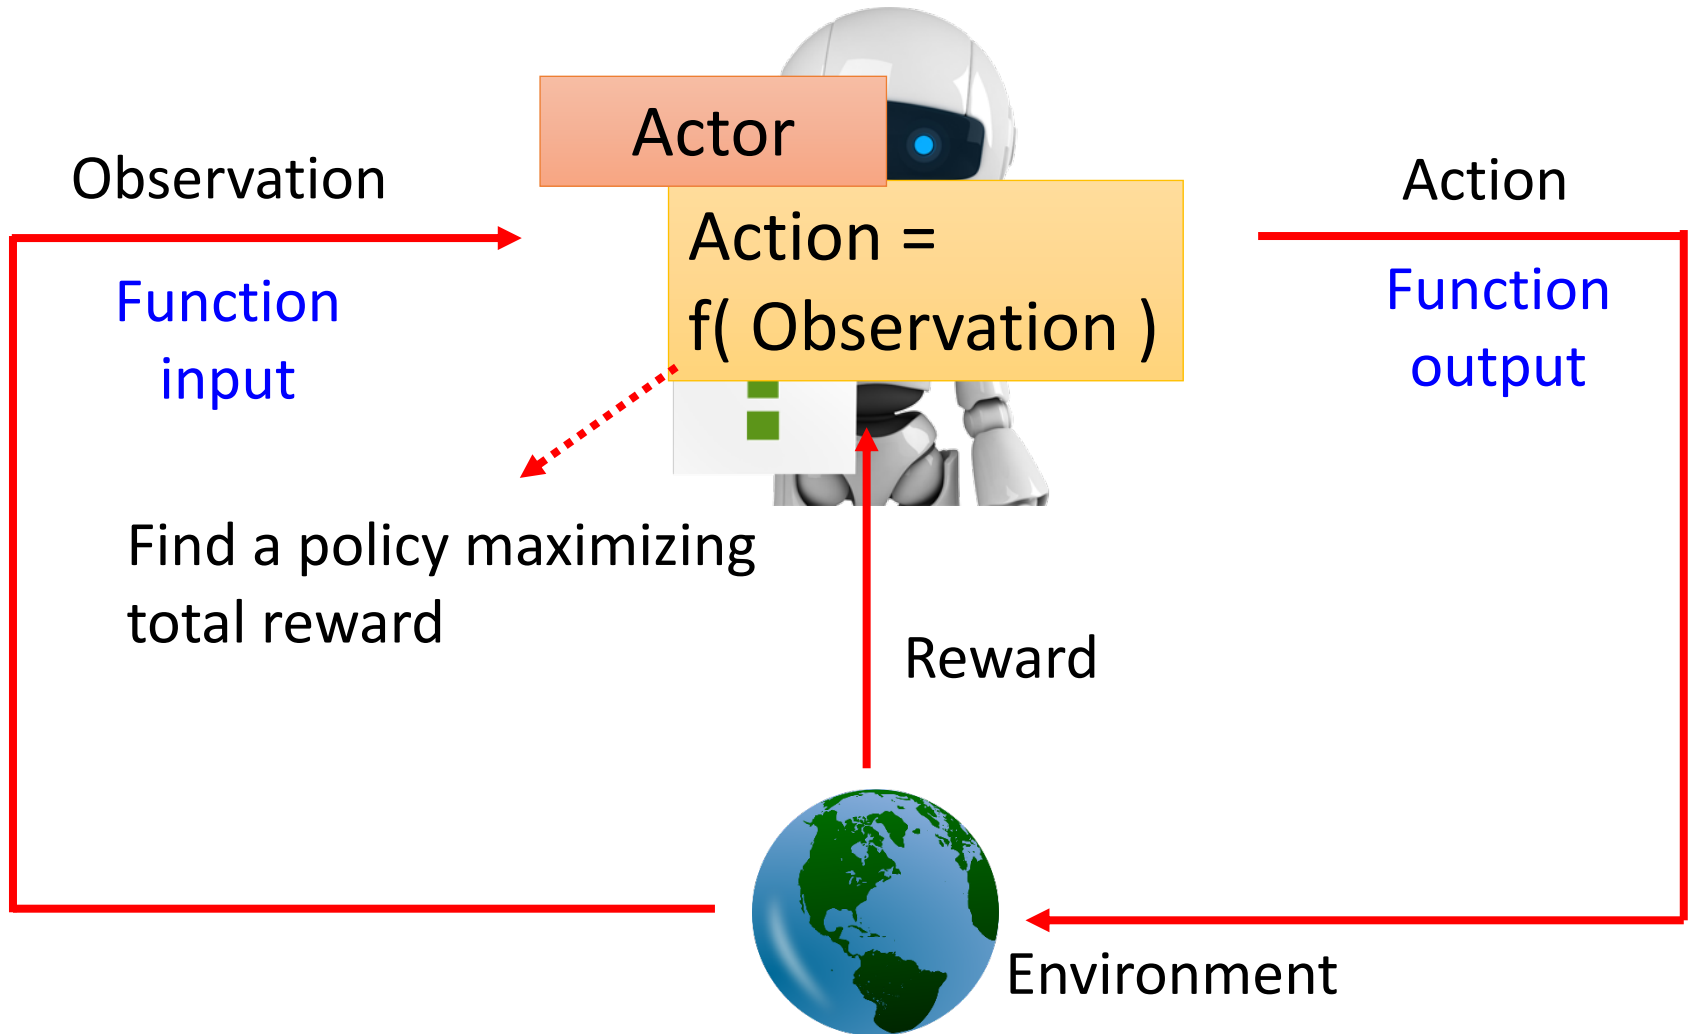
Outline

What is RL? (Three steps in ML)

Policy Gradient

Actor-Critic

Reward Shaping

No Reward: Learning from Demonstration

Machine Learning
≈ Looking for a Function

Observation

Actor

Action = f(Observation)

Function input

Action

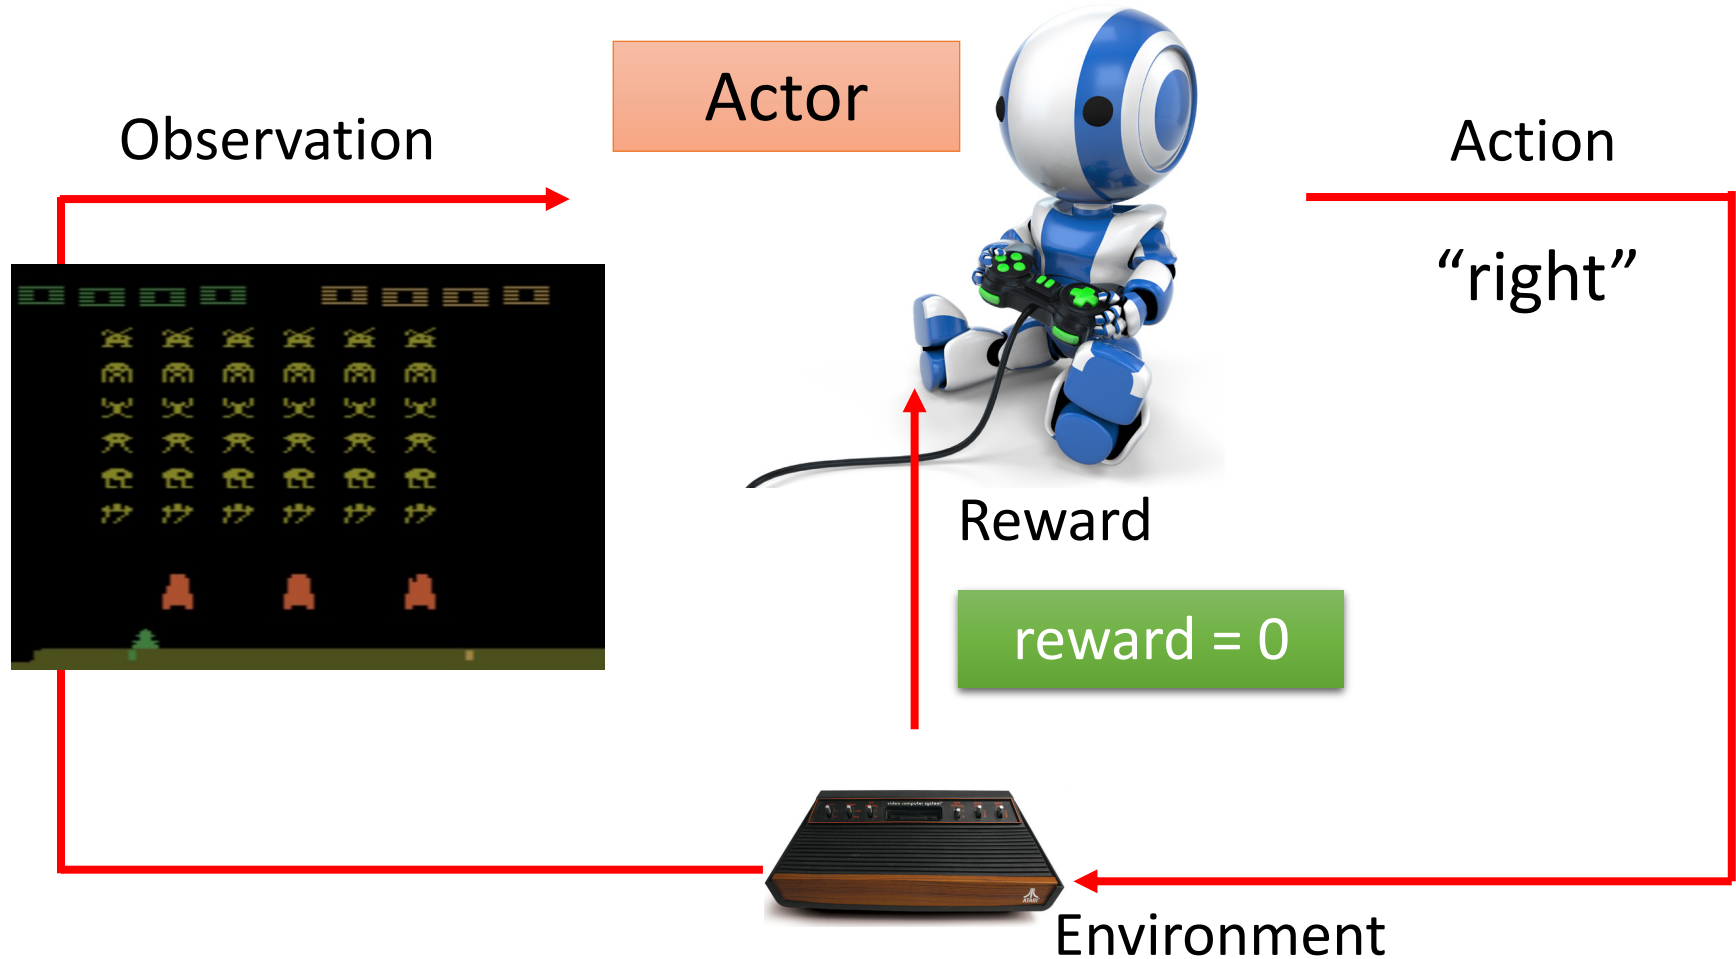Function output

Find a policy maximizing total reward

Reward

Environment

Example: Playing Video Game

- Space invader

Termination: all the aliens are killed, or your spaceship is destroyed.

Score (reward)

Kill the aliens

shield

fire

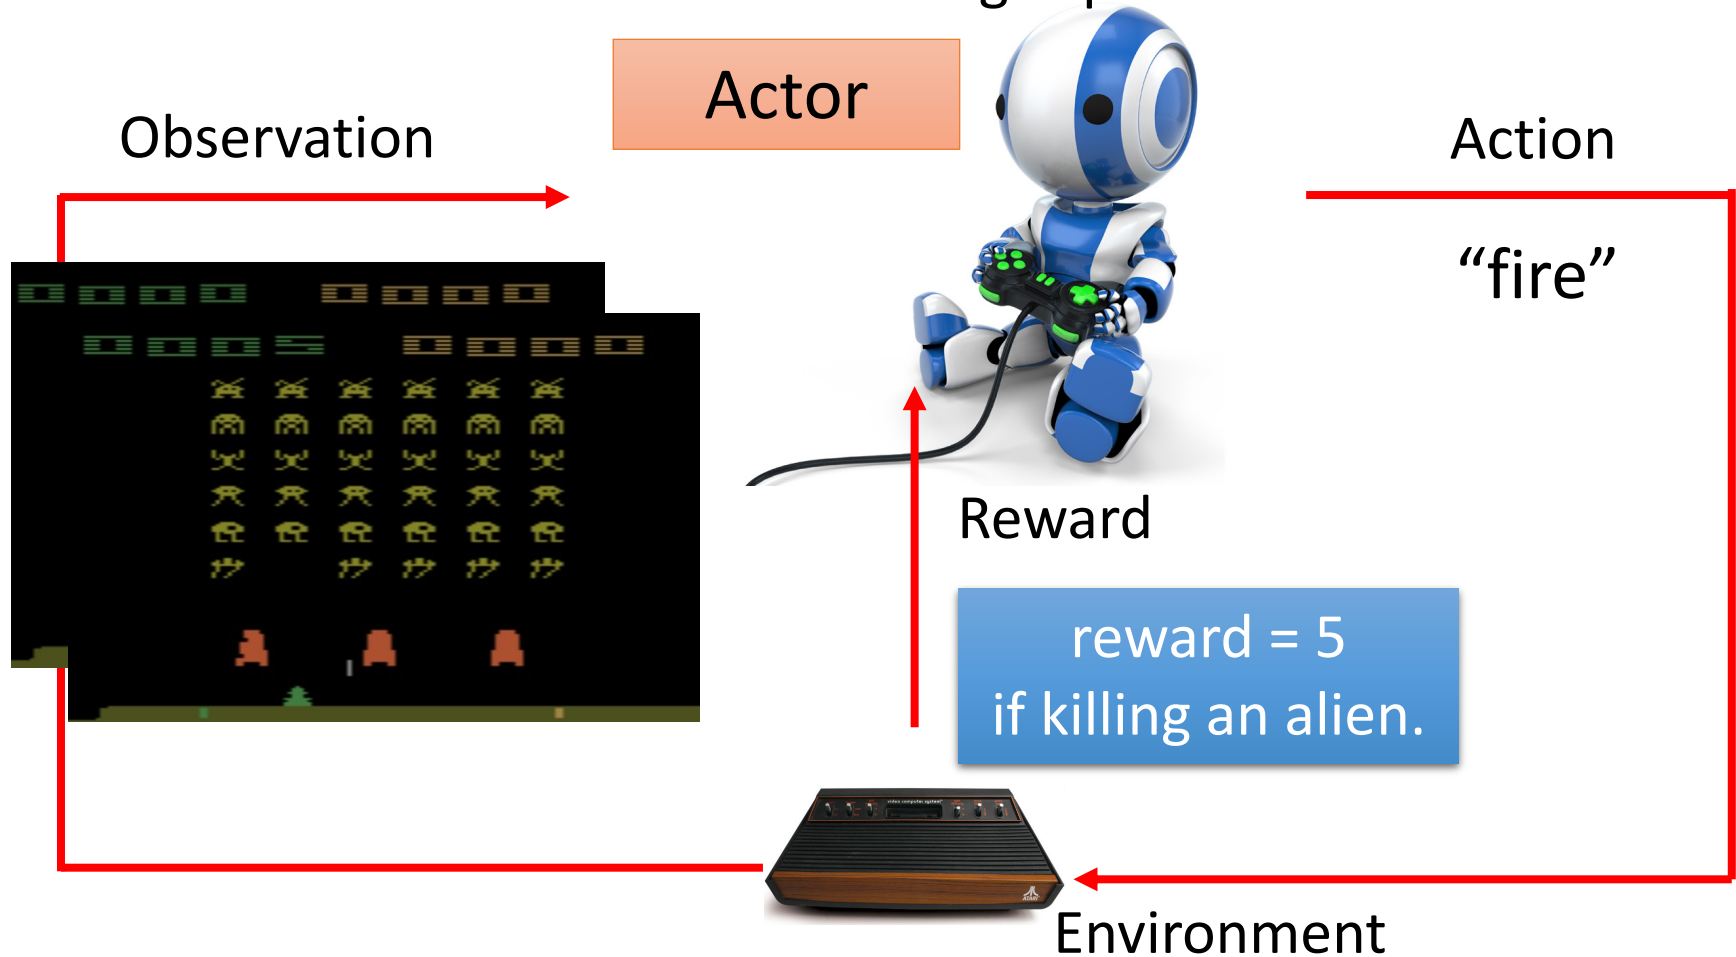# Example: Playing Video Game

Example: Playing Video Game

Find an actor maximizing expected reward.

Actor

Observation

Action

"fire"

Reward

reward = 5
if killing an alien.

Environment

Example: Learning to play Go

Observation

Action

AlphaGo

Reward

Next Move

Environment
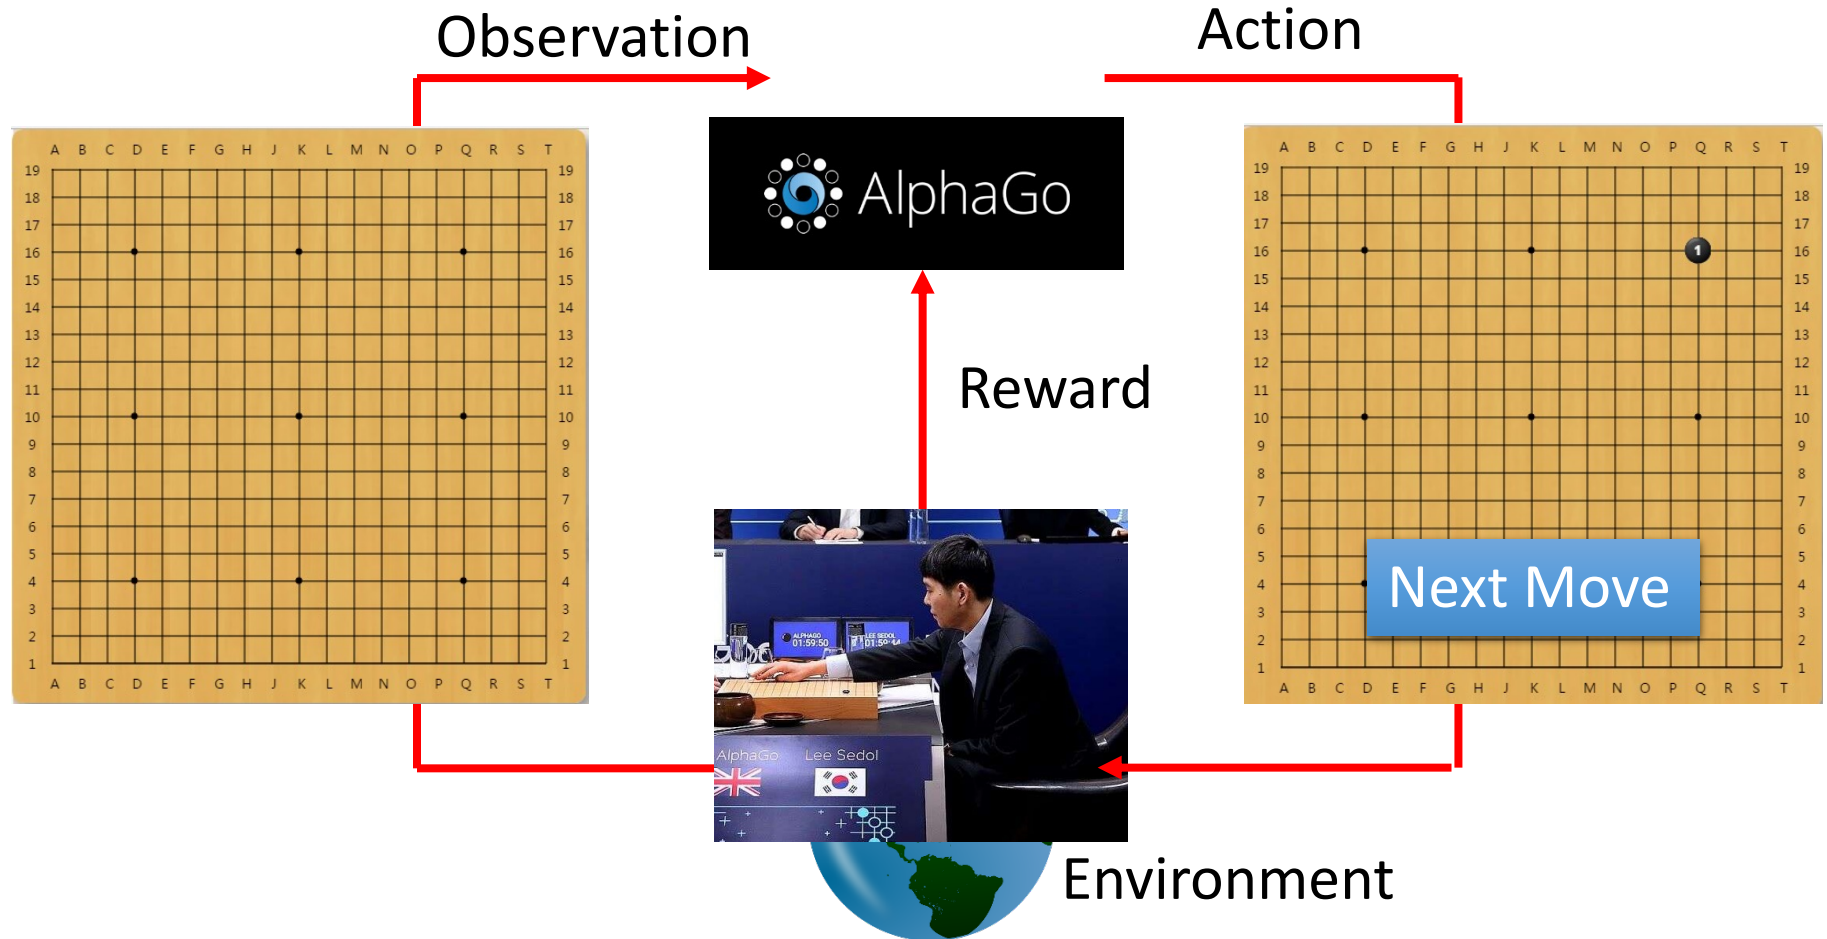
Example: Learning to play Go

Find an actor maximizing expected reward.

Observation

Action

AlphaGo

Reward

reward = 0 in most cases

If win, reward = 1

If loss, reward = -1

AlphaGo   Lee Sedol

Environment
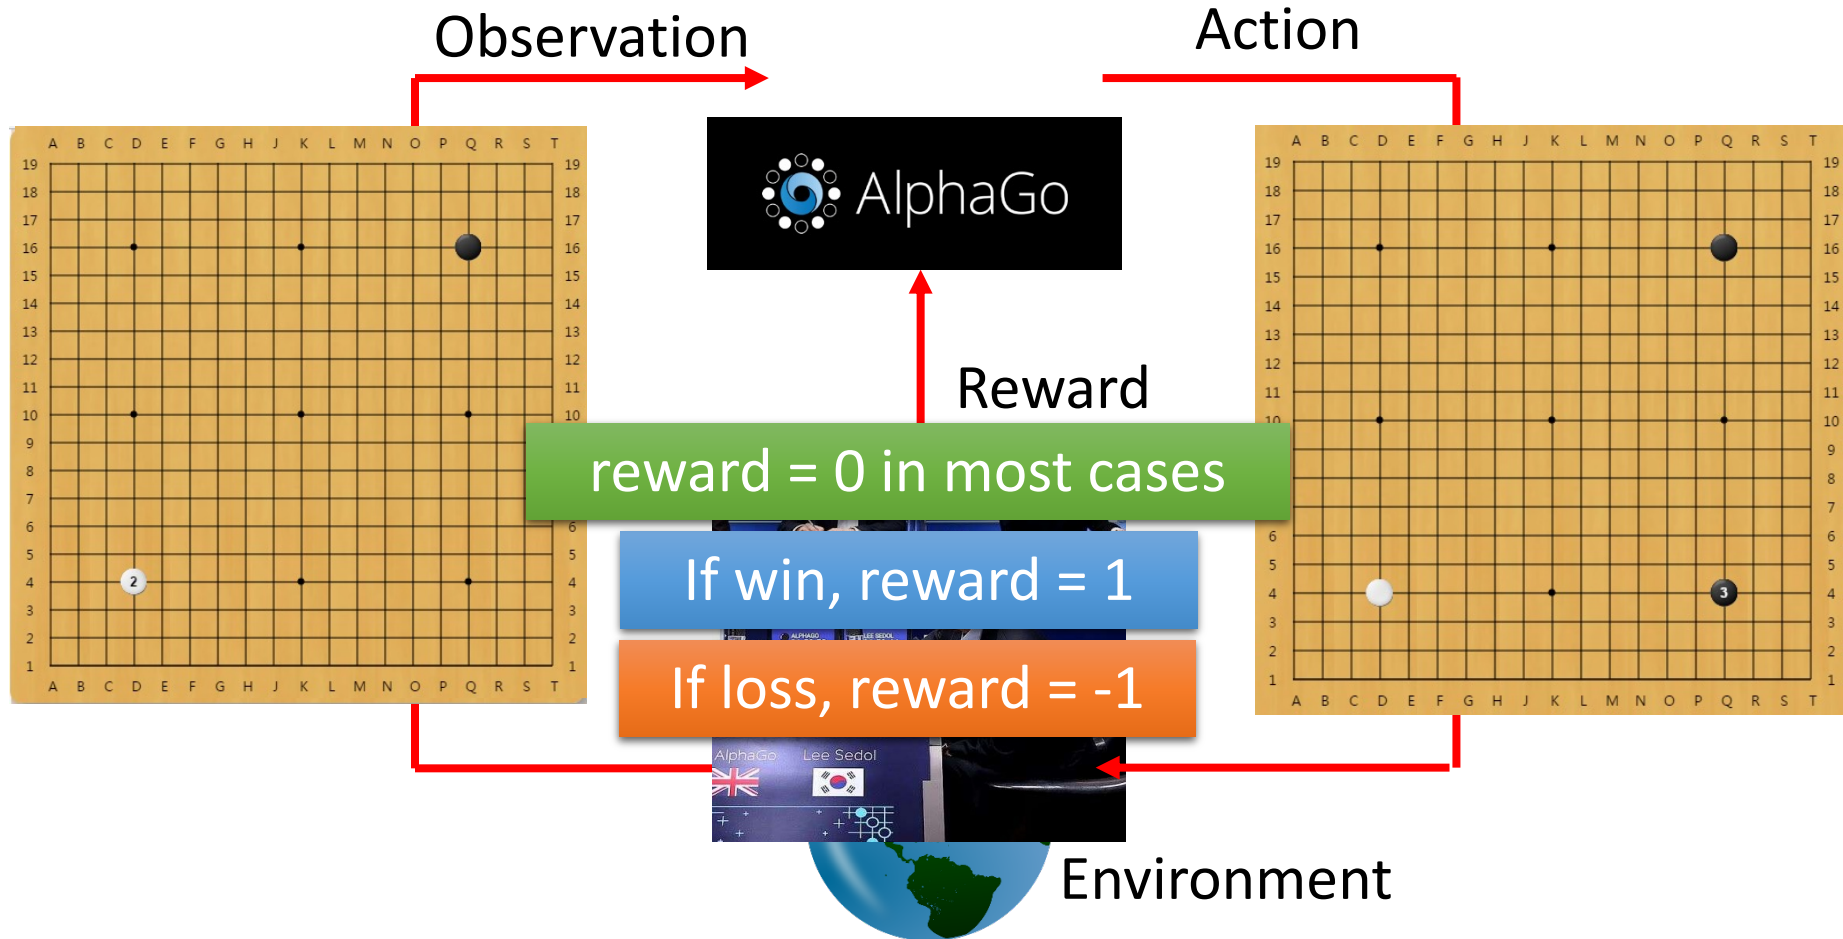
Machine Learning is so simple ……

Step 1: function with unknown variables → Step 2: define loss from training data → Step 3: optimization

Step 1: Function with Unknown

Policy Network (Actor)

pixels → [neural network] →

left 0.7
right 0.2
fire 0.1

Sample based on scores

Scores of actions

Classification Task!!!
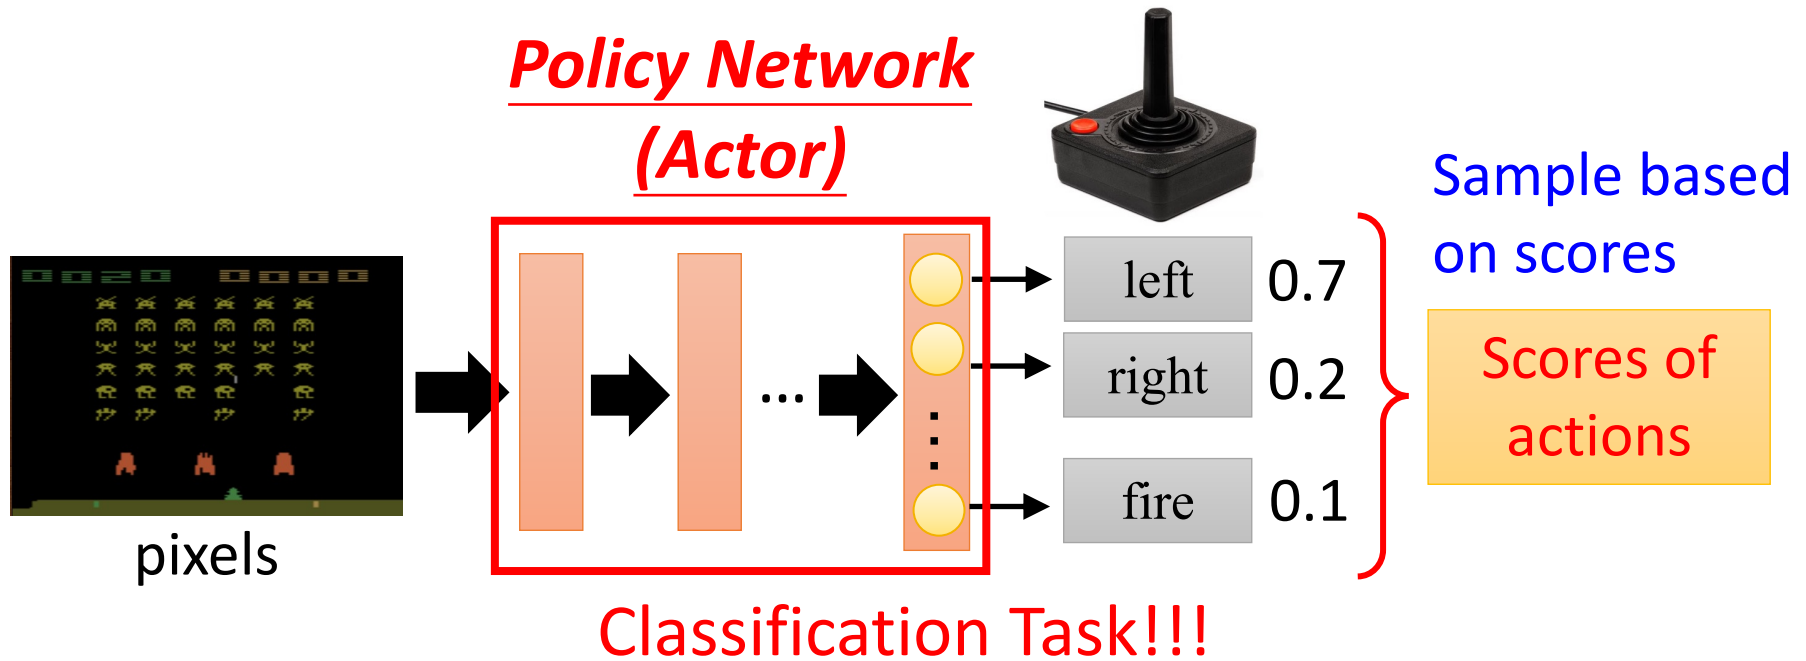
- Input of neural network: the observation of machine represented as a vector or a matrix
- Output neural network : each action corresponds to a neuron in output layer

Step 2: Define "Loss"

Start with
observation s_1

Observation s_2

Observation s_3

Obtain reward
$r_1 = 0$

Obtain reward
$r_2 = 5$

Action a_1: "right"

Action a_2 : "fire"

(kill an alien)

Step 2: Define "Loss"

Start with
observation s_1

Observation s_2

Observation s_3

This is an **_episode_**.

After many turns

.............▶

Game Over
(spaceship destroyed)

Total reward
(*return*):

$$R = \sum_{t=1}^{T} r_t$$

Obtain reward r_T

What we want
to maximize
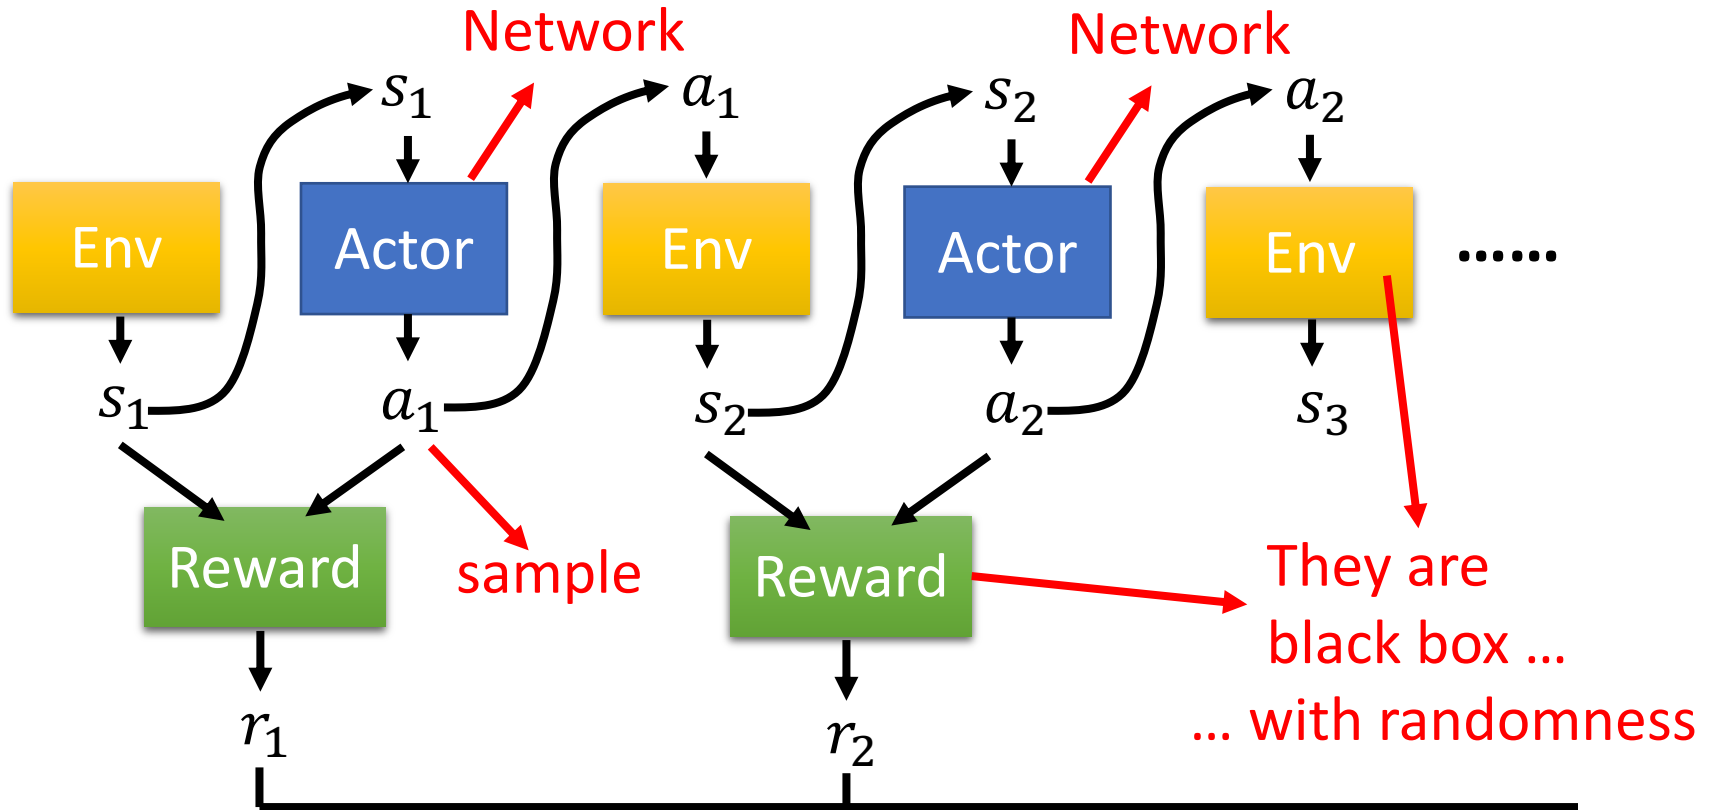
Action a_T

Outline

To learn more about policy gradient:
https://youtu.be/W8XF3ME8G2I

What is RL? (Three steps in ML)

Policy Gradient

Actor-Critic

Reward Shaping

No Reward: Learning from Demonstration

How to control your actor

- Make it take (or don't take) a specific action $\hat{a}$ given specific observation s.

Take action $\hat{a}$

$$L = e \text{ (loss)}$$

Don't take action $\hat{a}$

$$L = -e$$

$$\theta^* = arg\min_\theta L$$

How to control your actor

Take action $\hat{a}$ given s

a e_1 $\hat{a}$

left	1
right	0
fire	0

Don't take action $\hat{a}'$ given s'

a e_2 $\hat{a}'$

left	0
right	1
fire	0

$$L = e_1 - e_2 \qquad \theta^* = arg\min_\theta L$$

How to control your actor

Training Data

$\{s_1, \hat{a}_1\}$ +1 **Yes**

$\{s_2, \hat{a}_2\}$ -1 **No**

$\{s_3, \hat{a}_3\}$ +1 **Yes**

$\vdots$ $\vdots$

$\{s_N, \hat{a}_N\}$ -1 **No**

$s \longrightarrow$ Actor θ $\longrightarrow a$

$$L = + e_1 - e_2 + e_3 \cdots - e_N$$

$$\theta^* = arg \min_\theta L$$

How to control your actor

Training Data

$\{s_1, \hat{a}_1\}$ A_1 $+1.5$

$\{s_2, \hat{a}_2\}$ A_2 $- 0.5$

$\{s_3, \hat{a}_3\}$ A_3 $+0.5$

$\vdots$ $\vdots$

$\{s_N, \hat{a}_N\}$ A_N -10

?   ?

$$s \longrightarrow \boxed{\text{Actor} \quad \theta} \longrightarrow a$$

$$L = \sum A_n e_n$$

$$\theta^* = arg \min_{\theta} L$$

Version 0

Training Data

$\{s_1, a_1\}$ $A_1 = r_1$

$\{s_2, a_2\}$ $A_2 = r_2$

$\{s_3, a_3\}$ $A_3 = r_3$

$\vdots$ $\vdots$

$\{s_N, a_N\}$ $A_N = r_N$

many episodes

Short-sighted Version!

Version 0
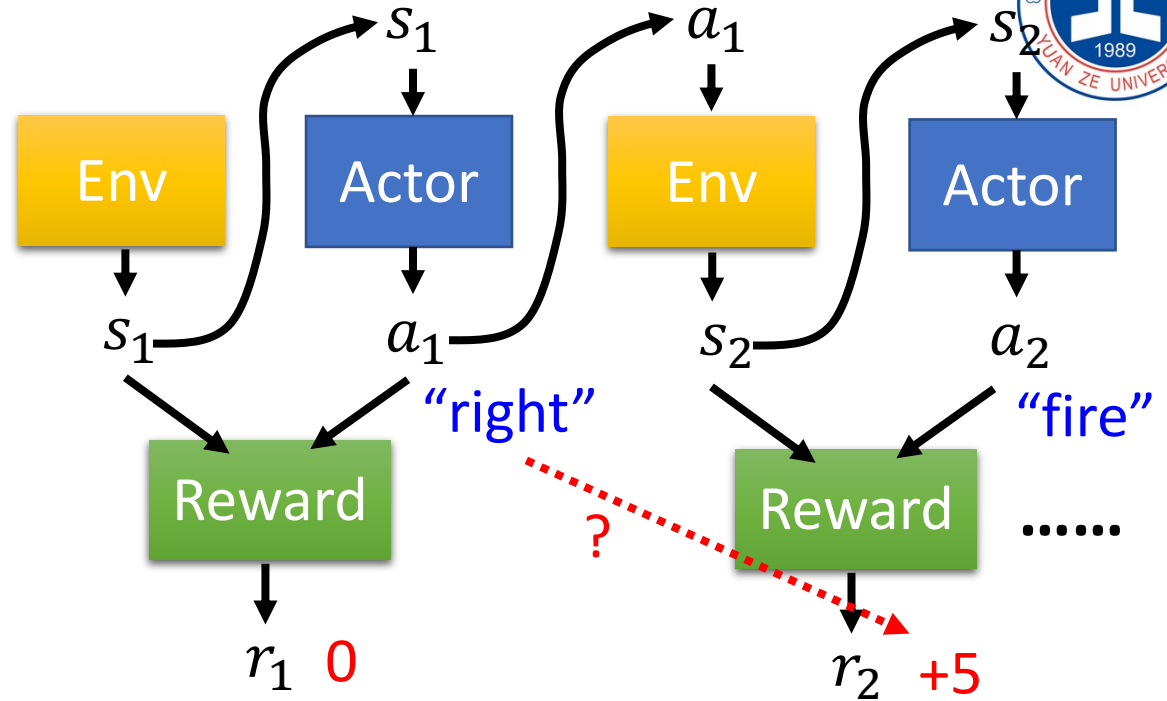
- An action affects the subsequent observations and thus subsequent rewards.

- *Reward delay*: Actor has to sacrifice immediate reward to gain more long-term reward.

- In space invader, only "fire" yields positive reward, so vision 0 will learn an actor that always "fire".

Version 1

s_1 $\qquad$ s_2 $\qquad$ s_3 $\qquad\qquad$ s_N

a_1 $\qquad$ a_2 $\qquad$ a_3 $\qquad$ ...... $\qquad$ a_N

r_1 $\qquad$ r_2 $\qquad$ r_3 $\qquad\qquad$ r_N

$$G_1 = r_1 + r_2 + r_3 + ... + r_N$$

$$G_2 = r_2 + r_3 + ... + r_N$$

$$G_3 = r_3 + ... + r_N$$

cumulated reward

Training Data

$\{s_1, a_1\}$ $\qquad$ $A_1 = G_1$

$\{s_2, a_2\}$ $\qquad$ $A_2 = G_2$

$\{s_3, a_3\}$ $\qquad$ $A_3 = G_3$

$\vdots$ $\qquad\qquad$ $\vdots$

$\{s_N, a_N\}$ $\qquad$ $A_N = G_N$

$$G_t = \sum_{n=t}^{N} r_n$$

Version 2

s_1 s_2 s_3 s_N

a_1 a_2 a_3 ...... a_N

r_1 r_2 r_3 r_N

Also the credit of a_1 ?

$$G_1 = r_1 + r_2 + r_3 + \ldots\ldots + \boxed{r_N}$$

$$G_1' = r_1 + \gamma r_2 + \gamma^2 r_3 + \ldots\ldots$$

Discount factor $\gamma < 1$

This is called "discounted cumulated reward."

$\{s_1, a_1\}$ $A_1 = G_1'$

$\{s_2, a_2\}$ $A_2 = G_2'$

$\{s_3, a_3\}$ $A_3 = G_3'$

$\vdots$ $\vdots$

$\{s_N, a_N\}$ $A_N = G_N'$

$$G_t' = \sum_{n=t}^{N} \gamma^{n-t} r_n$$

Version 3

s_1 s_2 s_3 $\qquad$ s_N

a_1 a_2 a_3 $\cdots\cdots$ a_N

r_1 r_2 r_3 r_N

Good or bad reward is "relative"

If all the $r_n \geq 10$

$\qquad r_n = 10$ is negative ...

Minus by a baseline b **???**

Make G'_t have positive and negative values

Training Data

$\{s_1, a_1\}$ $A_1 = G'_1 - b$

$\{s_2, a_2\}$ $A_2 = G'_2 - b$

$\{s_3, a_3\}$ $A_3 = G'_3 - b$

$\vdots \qquad \vdots$

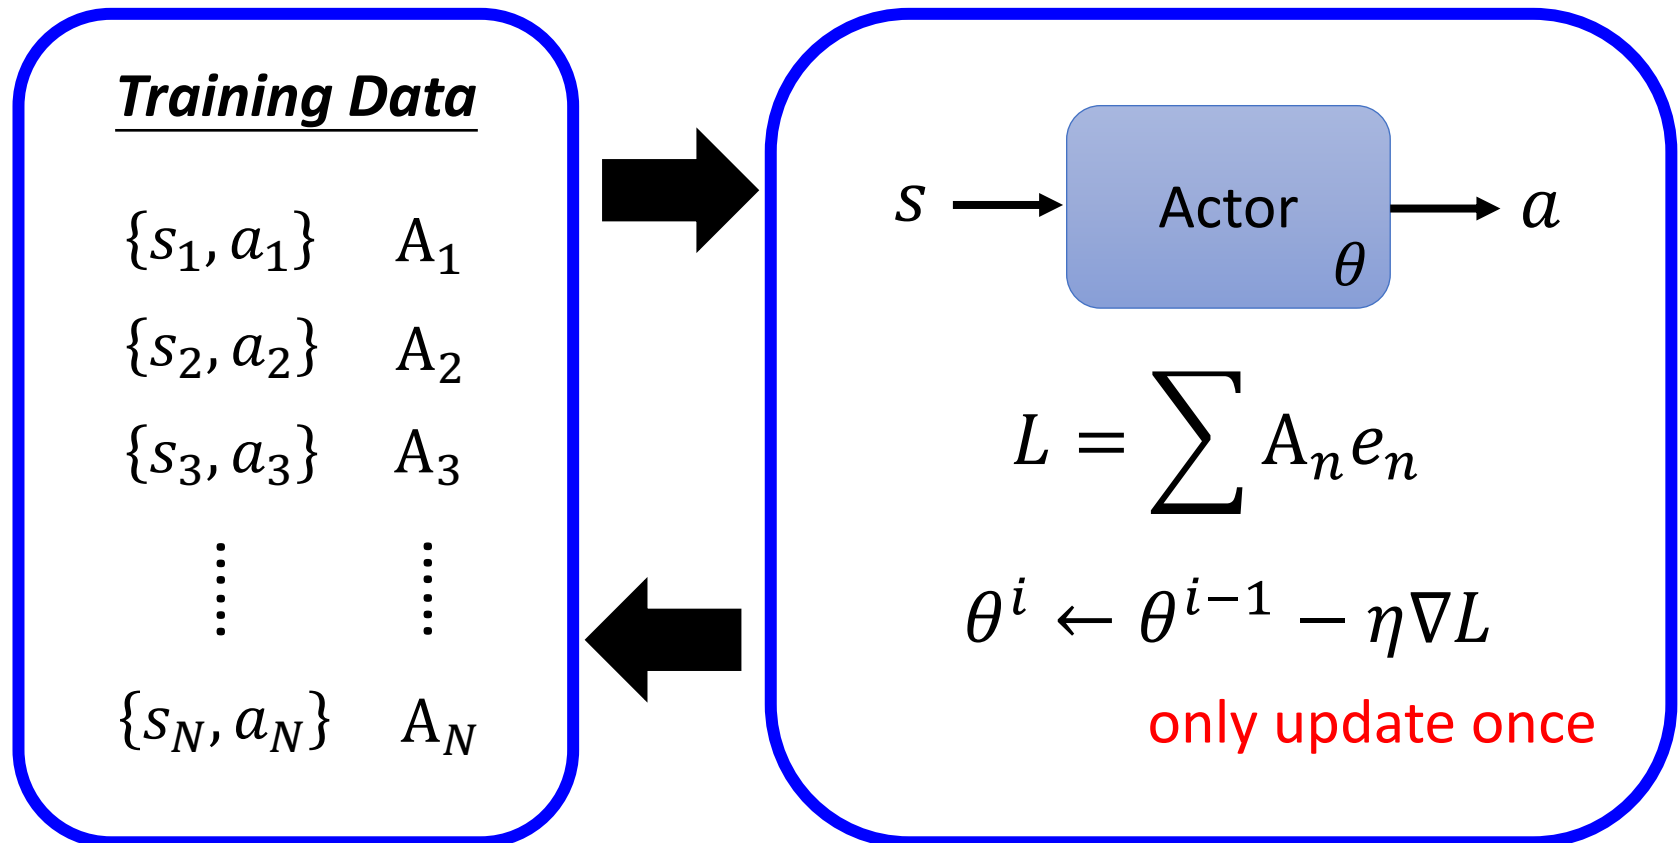$\{s_N, a_N\}$ $A_N = G'_N - b$

$$G'_t = \sum_{n=t}^{N} \gamma^{n-t} r_n$$

Policy Gradient

- Initialize actor network parameters θ^0

- For training iteration $i = 1$ to T
  - Using actor θ^{i-1} to interact
  - Obtain data $\{s_1, a_1\}, \{s_2, a_2\}, \dots, \{s_N, a_N\}$
  - Compute $A_1, A_2, \dots, A_N$
  - Compute loss L
  - $\theta^i \leftarrow \theta^{i-1} - \eta \nabla L$

Data collection is in the "for loop" of training iterations.

Policy Gradient

Training Data

$\{s_1, a_1\}$ A_1

$\{s_2, a_2\}$ A_2

$\{s_3, a_3\}$ A_3

$\vdots$ $\vdots$

$\{s_N, a_N\}$ A_N

$s \longrightarrow$ Actor $\longrightarrow a$

θ

$$L = \sum A_n e_n$$

$$\theta^i \leftarrow \theta^{i-1} - \eta \nabla L$$

only update once

Each time you update the model parameters, you need to collect the whole training set again.

Policy Gradient

- Initialize actor network parameters θ^0

- For training iteration $i = 1$ to T

  - Using actor $\boxed{\theta^{i-1}}$ to interact
  - Obtain data $\{s_1, a_1\}, \{s_2, a_2\}, \ldots, \{s_N, a_N\}$
  - Compute $A_1, A_2, \ldots, A_N$
  - Compute loss L
  - $\boxed{\theta^i} \leftarrow \boxed{\theta^{i-1}} - \eta \nabla L$

Experience of θ^{i-1}

May not be good for θ^i

One man's meat is another man's poison.

棋魂第八集

棋魂第八集

相對地，大馬步飛之後的下法會很複雜，出錯率也相對地提高。

要是大馬步飛有100手的話，小馬步飛就只有99手。

之前走小馬步飛是對的。

因為小馬步飛的後續比較容易推測，也不容易出錯。

Policy Gradient

- Initialize actor network parameters θ^0
- For training iteration $i = 1$ to T
    - Using actor θ^{i-1} to interact
    - Obtain data $\{s_1, a_1\}, \{s_2, a_2\}, \ldots, \{s_N, a_N\}$
    - Compute $A_1, A_2, \ldots, A_N$
    - Compute loss L
    - $\theta^i \leftarrow \theta^{i-1} - \eta \nabla L$

May not observe by θ^i

Trajectory of θ^{i-1}

s_1	s_2	s_3		s_N
a_1	a_2	a_3	$\cdots\cdots$	a_N
r_1	r_2	r_3		r_N

On-policy v.s. Off-policy

- The actor to train and the actor for interacting is the same. → On-policy

- Can the actor to train and the actor for interacting be different? → Off-policy

Trajectory of θ^{i-1}

$$
\begin{array}{cccccc}
s_1 & s_2 & s_3 & & s_N \\
a_1 & a_2 & a_3 & \cdots\cdots & a_N \\
r_1 & r_2 & r_3 & & r_N
\end{array}
$$

→ To train θ^i

In this way, we do not have to collection data after each update.

# Off-policy → Proximal Policy Optimization (PPO)

- The actor to train has to know its difference from the actor to interact.

video:

https://youtu.be/OAKAZhFmYoI

Not apply to everyone

the actor to train
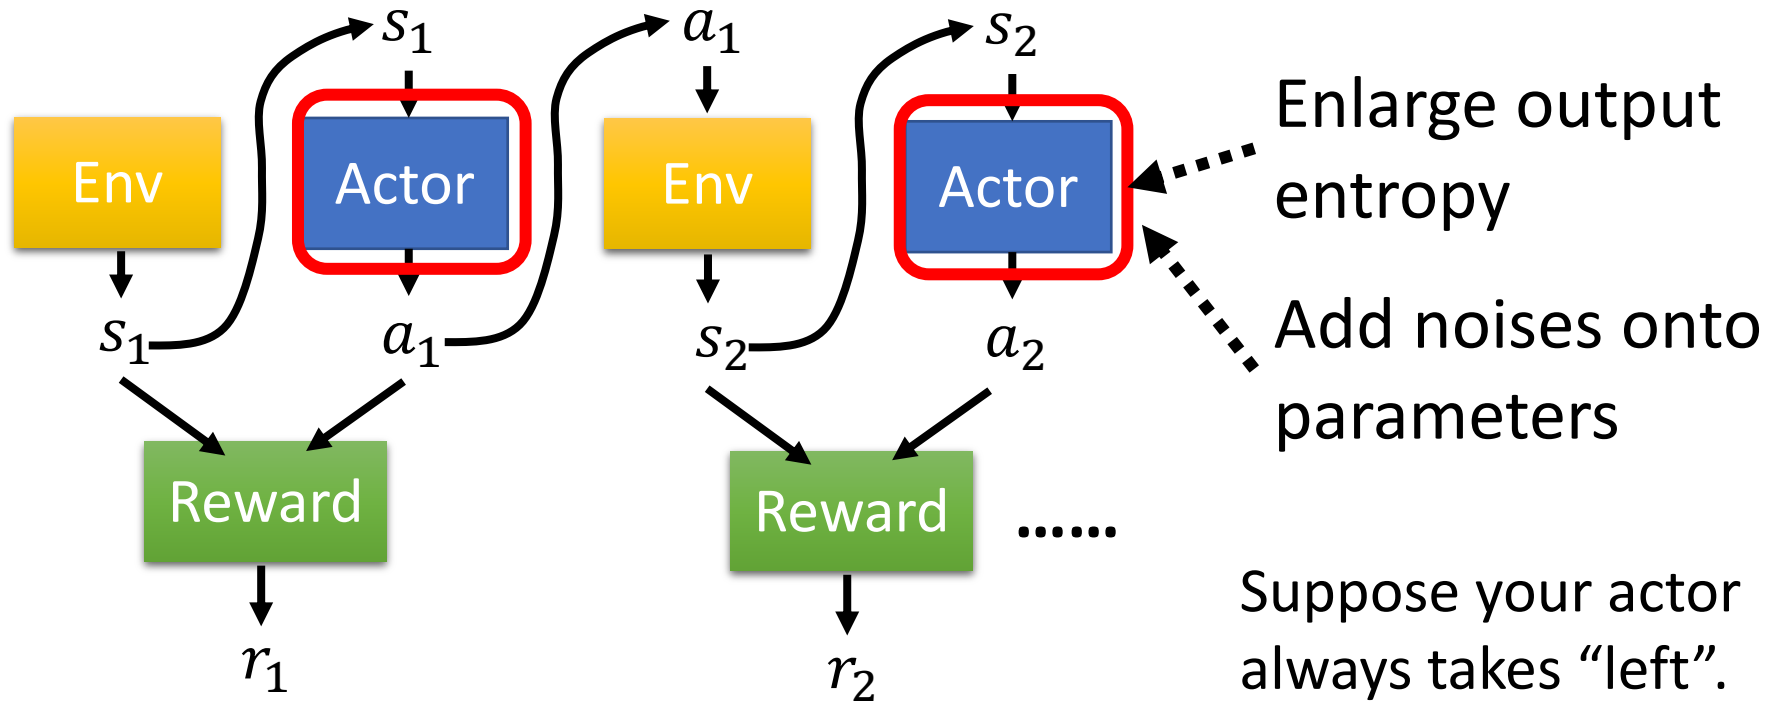
https://disp.cc/b/115-bLHe

如果喜歡一個女孩
就勇敢去告白
我用這個方法
從來沒失敗過

——克里斯·伊凡

the actor to interact

Collection Training Data: Exploration

The actor needs to have randomness during data collection.

A major reason why we sample actions. ☺

Enlarge output entropy

Add noises onto parameters

Suppose your actor always takes "left".

We never know what would happen if taking "fire".

DeepMind - PPO

OpenAI – PPO

Outline

What is RL? (Three steps in ML)

Policy Gradient

Actor-Critic

Reward Shaping

No Reward: Learning from Demonstration

Critic

$$G_1' = r_1 + \gamma r_2 + \gamma^2 r_3 + \ldots\ldots$$

- Critic: Given actor θ, how good it is when observing s (and taking action a)

- Value function $V^\theta(s)$: When using actor θ, the discounted *cumulated* reward expects to be obtained after seeing s

$V^\theta(s)$ is large $V^\theta(s)$ is smaller

The output values of a critic depend on the actor evaluated.

How to estimate $V^\theta(s)$

- **Monte-Carlo (MC) based approach**

  The critic watches actor θ to interact with the environment.

After seeing s_a,

  Until the end of the episode, the discounted cumulated reward is G'_a

After seeing s_b,

  Until the end of the episode, the discounted cumulated reward is G'_b

$$s_a \rightarrow V^\theta \rightarrow V^\theta(s_a) \leftrightarrow G'_a$$

$$s_b \rightarrow V^\theta \rightarrow V^\theta(s_b) \leftrightarrow G'_b$$

How to estimate $V^\pi(s)$
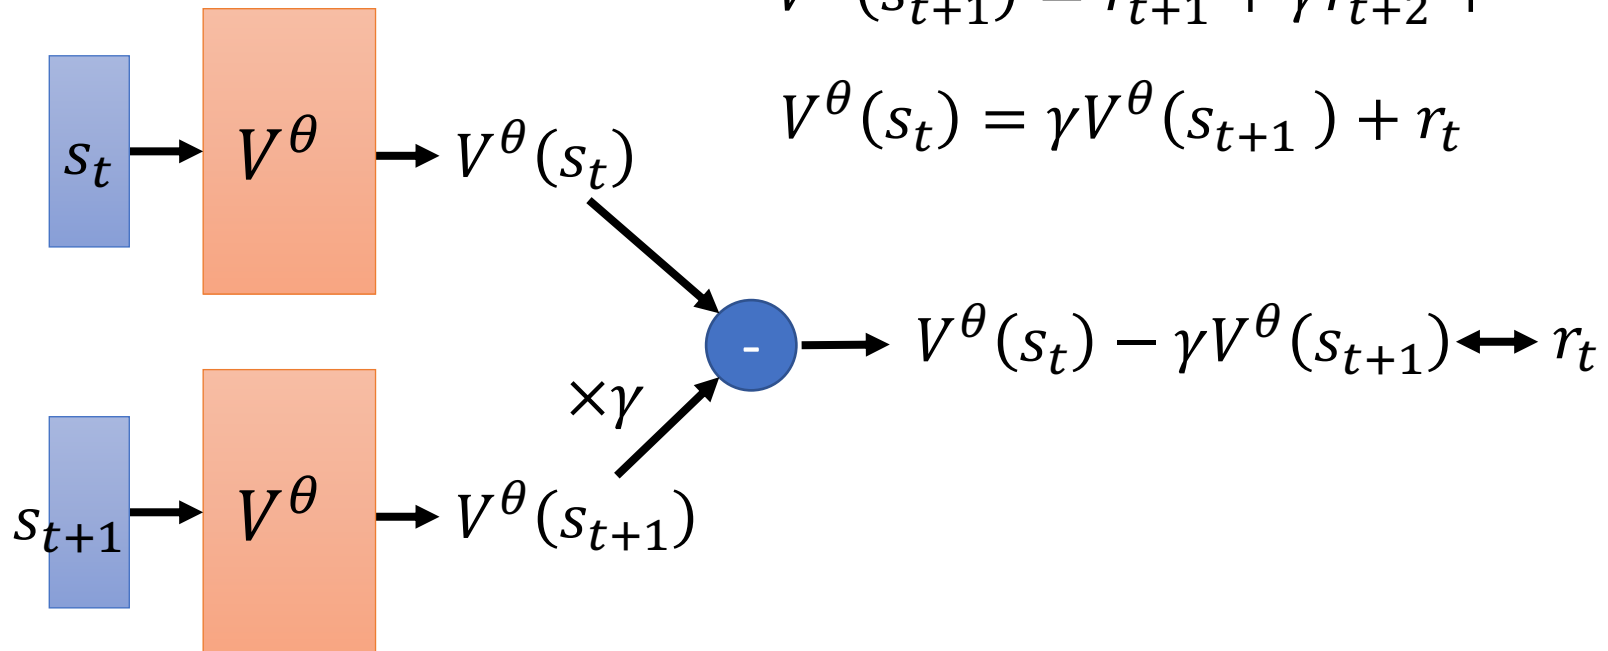
- **Temporal-difference (TD) approach**

$$\cdots s_t, a_t, r_t, s_{t+1} \cdots$$

(ignore the expectation here)

$$V^\theta(s_t) = r_t + \gamma r_{t+1} + \gamma^2 r_{t+2} \ldots$$

$$V^\theta(s_{t+1}) = r_{t+1} + \gamma r_{t+2} + \cdots$$

$$V^\theta(s_t) = \gamma V^\theta(s_{t+1}) + r_t$$

$$V^\theta(s_t) - \gamma V^\theta(s_{t+1}) \longleftrightarrow r_t$$

MC v.s. TD

- The critic has observed the following 8 episodes
  - $s_a, r = 0, s_b, r = 0$, END
  - $s_b, r = 1$, END
  - $s_b, r = 1$, END
  - $s_b, r = 1$, END
  - $s_b, r = 1$, END
  - $s_b, r = 1$, END
  - $s_b, r = 1$, END
  - $s_b, r = 0$, END

(Assume γ = 1, and the actions are ignored here.)

$$V^\theta(s_b) = 3/4$$

$$V^\theta(s_a) = ? \quad \textcolor{red}{0?} \quad \textcolor{red}{3/4?}$$

Monte-Carlo: $V^\theta(s_a) = 0$

Temporal-difference:

$$V^\theta(s_a) = V^\theta(s_b) + r$$

$$\textcolor{red}{3/4} \quad\quad \textcolor{red}{3/4} \quad\quad \textcolor{red}{0}$$

Version 3.5

$$s_1 \quad s_2 \quad s_3 \qquad\qquad s_N$$

$$a_1 \quad a_2 \quad a_3 \quad \cdots\cdots \quad a_N$$

$$r_1 \quad r_2 \quad r_3 \qquad\qquad r_N$$

$$s \rightarrow V^\theta \rightarrow V^\theta(s)$$

Training Data

$$\{s_1, a_1\} \quad \mathrm{A}_1 = G'_1 - b$$

$$\{s_2, a_2\} \quad \mathrm{A}_2 = G'_2 - b$$

$$\{s_3, a_3\} \quad \mathrm{A}_3 = G'_3 - b$$

$$\vdots \qquad\qquad \vdots$$

$$\{s_N, a_N\} \quad \mathrm{A}_N = G'_N - b$$

Version 3.5

$$s_1 \quad s_2 \quad s_3 \qquad\qquad s_N$$

$$a_1 \quad a_2 \quad a_3 \quad \cdots\cdots \quad a_N$$

$$r_1 \quad r_2 \quad r_3 \qquad\qquad r_N$$

$s \rightarrow V^\theta \rightarrow V^\theta(s)$

Training Data

$$\{s_1, a_1\} \quad A_1 = G_1' - V^\theta(s_1)$$

$$\{s_2, a_2\} \quad A_2 = G_2' - V^\theta(s_2)$$

$$\{s_3, a_3\} \quad A_3 = G_3' - V^\theta(s_3)$$

$$\vdots \qquad\qquad \vdots$$

$$\{s_N, a_N\} \quad A_N = G_N' - V^\theta(s_N)$$

Version 3.5

$\{s_t, a_t\}$ $A_t = G_t' - V^\theta(s_t)$


$G = 100$

$G = 3$

$G = 1$

$G = 2$

$G = -10$

$V^\theta(s_t)$

s_t

(not necessary take a_t)

(You sample the actions based on a distribution)

$A_t > 0$

a_t is better than average.

a_t
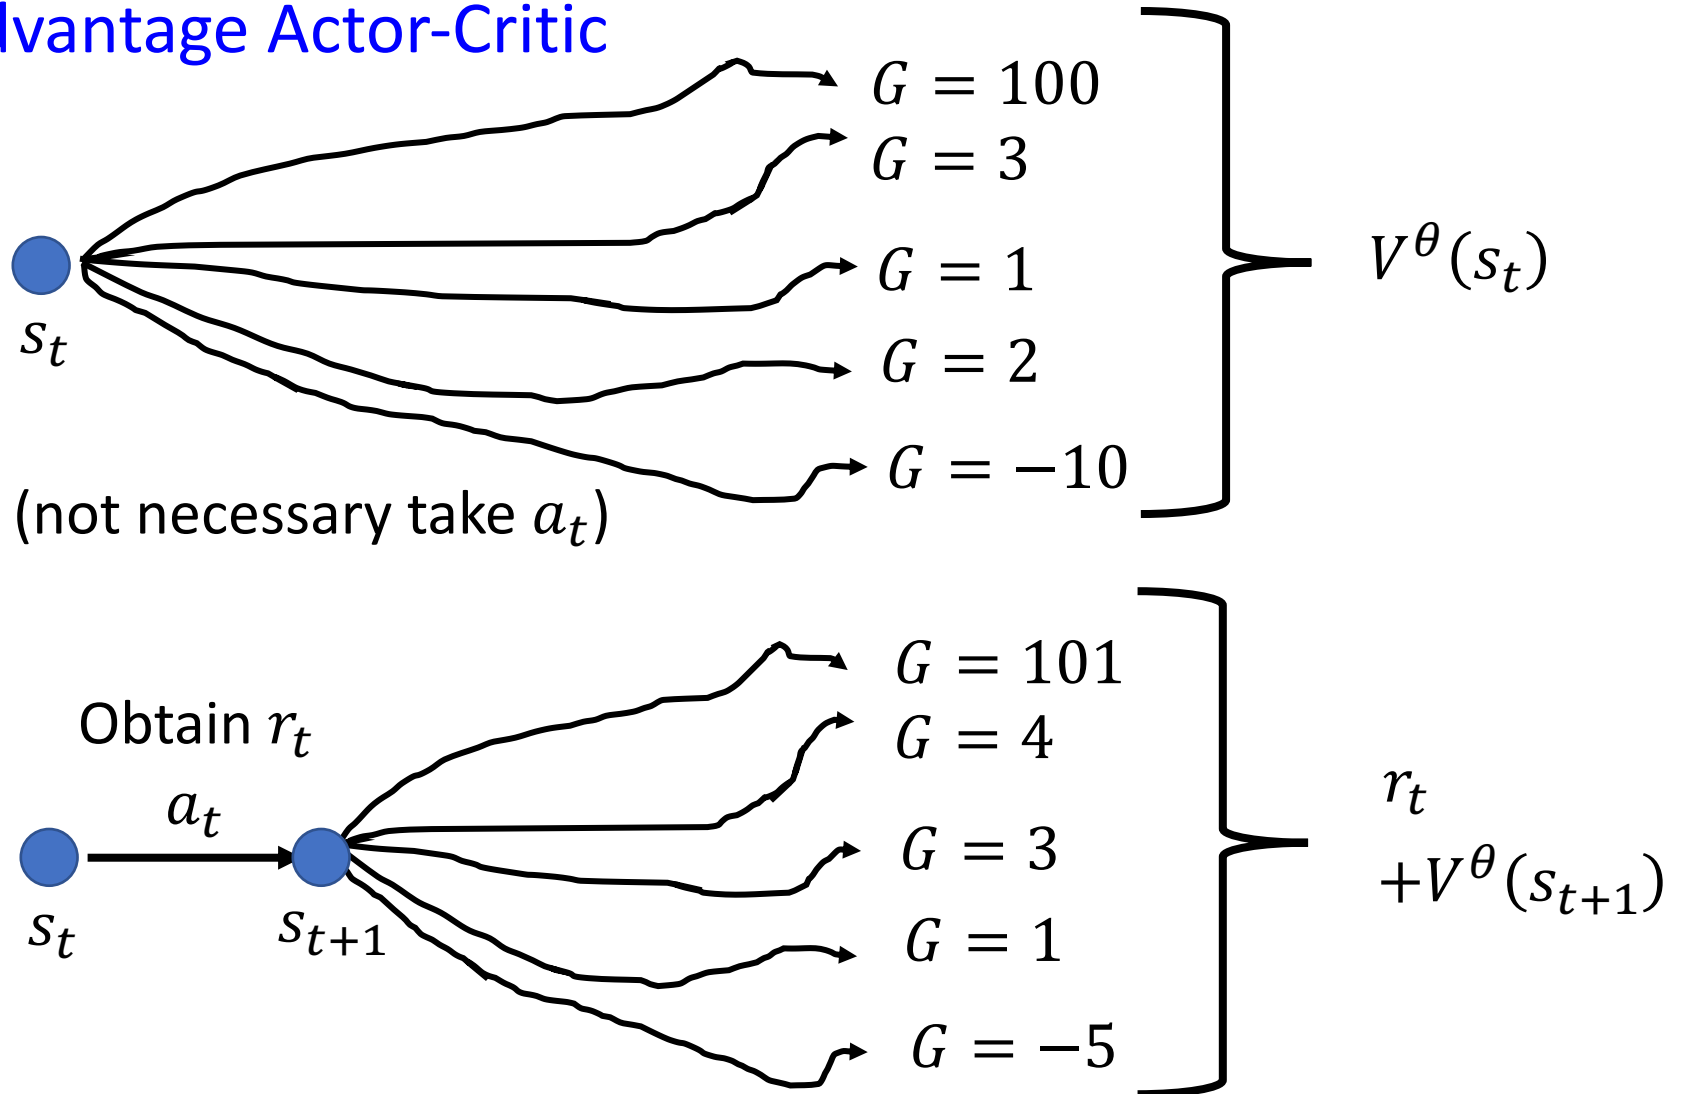
s_t

G_t' $A_t < 0$

a_t is worse than average.

Just a sample

Version 4

$$r_t + V^\theta(s_{t+1}) - V^\theta(s_t)$$

$$\{s_t, a_t\} \quad A_t = \cancel{G_t' - V^\theta(s_t)}$$

Advantage Actor-Critic

$G = 100$

$G = 3$

$G = 1$

$G = 2$

$G = -10$

$V^\theta(s_t)$

(not necessary take a_t)

Obtain r_t

a_t

s_t s_{t+1}

$G = 101$

$G = 4$

$G = 3$

$G = 1$

$G = -5$

r_t
$+V^\theta(s_{t+1})$

Tip of Actor-Critic

- The parameters of actor and critic can be shared.

Outlook: Deep Q Network (DQN)

Video:
https://youtu.be/o_g9JUMw1Oc
https://youtu.be/2-zGCx4iv_k

https://arxiv.org/abs/1710.02298

Outline

What is RL? (Three steps in ML)

Policy Gradient

Actor-Critic

Reward Shaping

No Reward: Learning from Demonstration

Sparse Reward

$$A_t = r_t + V^\theta(s_{t+1}) - V^\theta(s_t)$$

Training Data

$\{s_1, a_1\}$ A_1

$\{s_2, a_2\}$ A_2

$\{s_3, a_3\}$ A_3

$\vdots$ $\vdots$

$\{s_N, a_N\}$ A_N

$$
\begin{array}{cccccc}
s_1 & s_2 & s_3 & & & s_N \\
a_1 & a_2 & a_3 & \cdots\cdots & & a_N \\
r_1 & r_2 & r_3 & & & r_N
\end{array}
$$

If $r_t = 0$ in most cases ➡️ We don't know actions are good or bad.

e.g., robot arm to bolt on the screws

The developers define extra rewards to guide agents. ➡️ **_reward shaping_**

Reward Shaping

VizDoom   https://openreview.net/forum?id=Hk3mPK5gg&noteId=Hk3mPK5gg

Visual Doom AI Competition @ CIG 2016
https://www.youtube.com/watch?v=94EPSjQH38Y

Reward Shaping

Parameters	Description	FlatMap	CIGTrack1
living	Penalize agent who just lives	-0.008 / action	
health_loss	Penalize health decrement	-0.05 / unit	
ammo_loss	Penalize ammunition decrement	-0.04 / unit	
health_pickup	Reward for medkit pickup	0.04 / unit	
ammo_pickup	Reward for ammunition pickup	0.15 / unit	
dist_penalty	Penalize the agent when it stays	-0.03 / action	
dist_reward	Reward the agent when it moves	9e-5 / unit distance	

https://bair.berkeley.edu/blog/2017/12/20/reverse-curriculum/

Reward Shaping - Curiosity

Obtaining extra reward when the agent sees something new (but meaningful).

Curiosity Driven Exploration by Self-Supervised Prediction

ICML 2017

Deepak Pathak, Pulkit Agrawal, Alexei Efros, Trevor Darrell
UC Berkeley

Source of video: https://pathak22.github.io/noreward-rl/

Outline

What is RL? (Three steps in ML)

Policy Gradient

Actor-Critic

Reward Shaping

No Reward: Learning from Demonstration

Motivation

- Even define reward can be challenging in some tasks.

- Hand-crafted rewards can lead to uncontrolled behavior.

Three Laws of Robotics:
1. A robot may not injure a human being or, through inaction, allow a human being to come to harm.
2. A robot must obey the orders given it by human beings except where such orders would conflict with the First Law.
3. A robot must protect its own existence as long as such protection does not conflict with the First or Second Laws.

?????

restraining individual human behavior and sacrificing some humans will ensure humanity's survival

Imitation Learning

Actor can interact with the environment, but reward function is not available

We have demonstration of the expert.

$$\{\hat{\tau}_1, \hat{\tau}_2, \cdots, \hat{\tau}_K\}$$

Each $\hat{\tau}$ is a trajectory of the export.

Self driving: record human drivers

Robot: grab the arm of robot

# Isn't it Supervised Learning?

- Self-driving cars as example

$$\hat{\tau} = \{s_1, \hat{a}_1, s_2, \hat{a}_2, \cdots\}$$

forward

$s_i \longrightarrow$ $\longrightarrow a_i \longleftrightarrow \hat{a}_i$

Problem: The experts only sample limited observation.

Yes, also known as *Behavior Cloning*

?

Expert

More problem ......

The agent will copy every behavior, even irrelevant actions.

https://www.youtube.com/watch?v=j2FSB3bseek

Inverse Reinforcement Learning

demonstration
of the expert

Environment

$\{\hat{\tau}_1, \hat{\tau}_2, \cdots, \hat{\tau}_K\}$

Reward
Function

Reinforcement
Learning

Expert

Using the reward function to find the
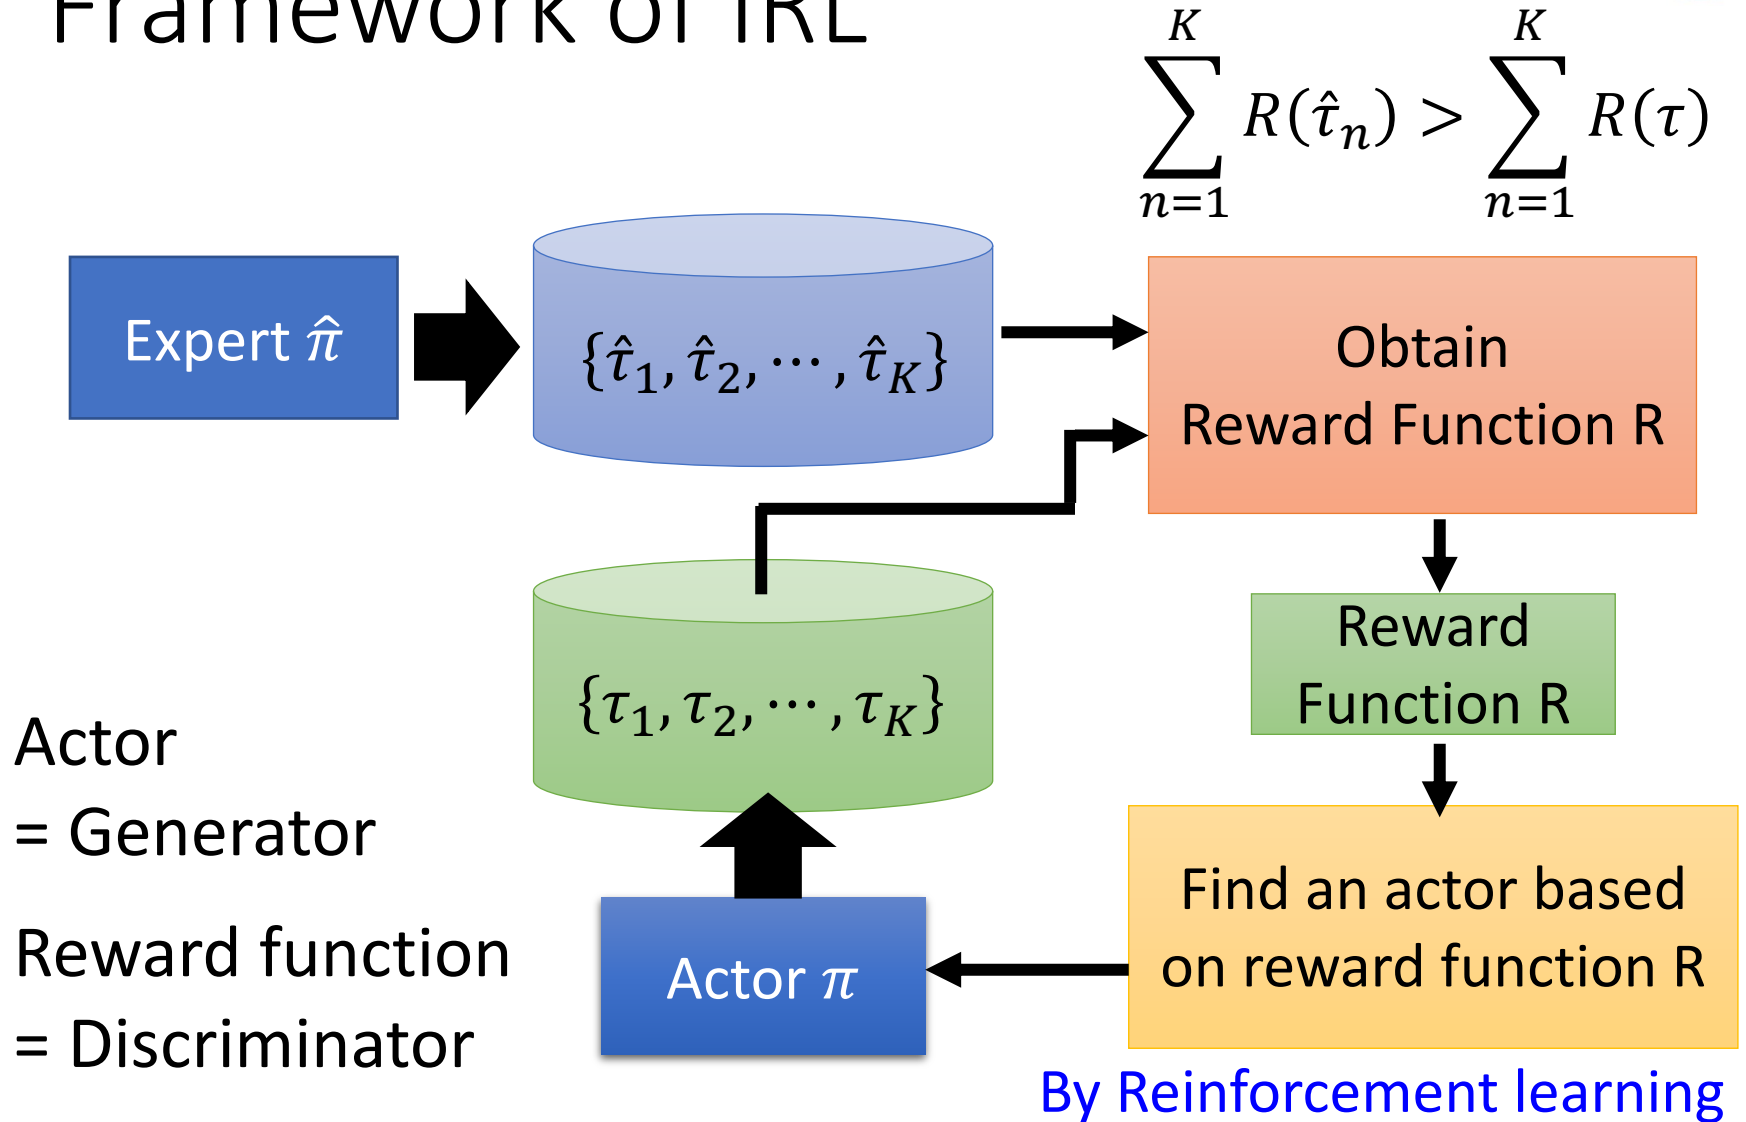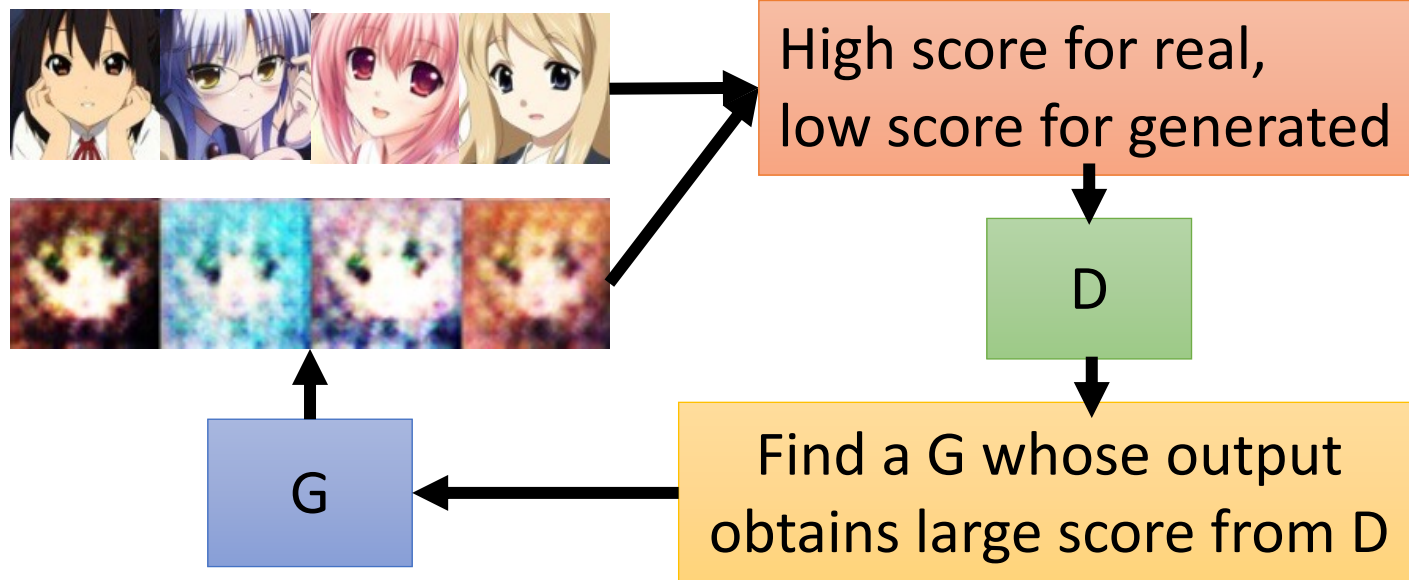***optimal actor***.
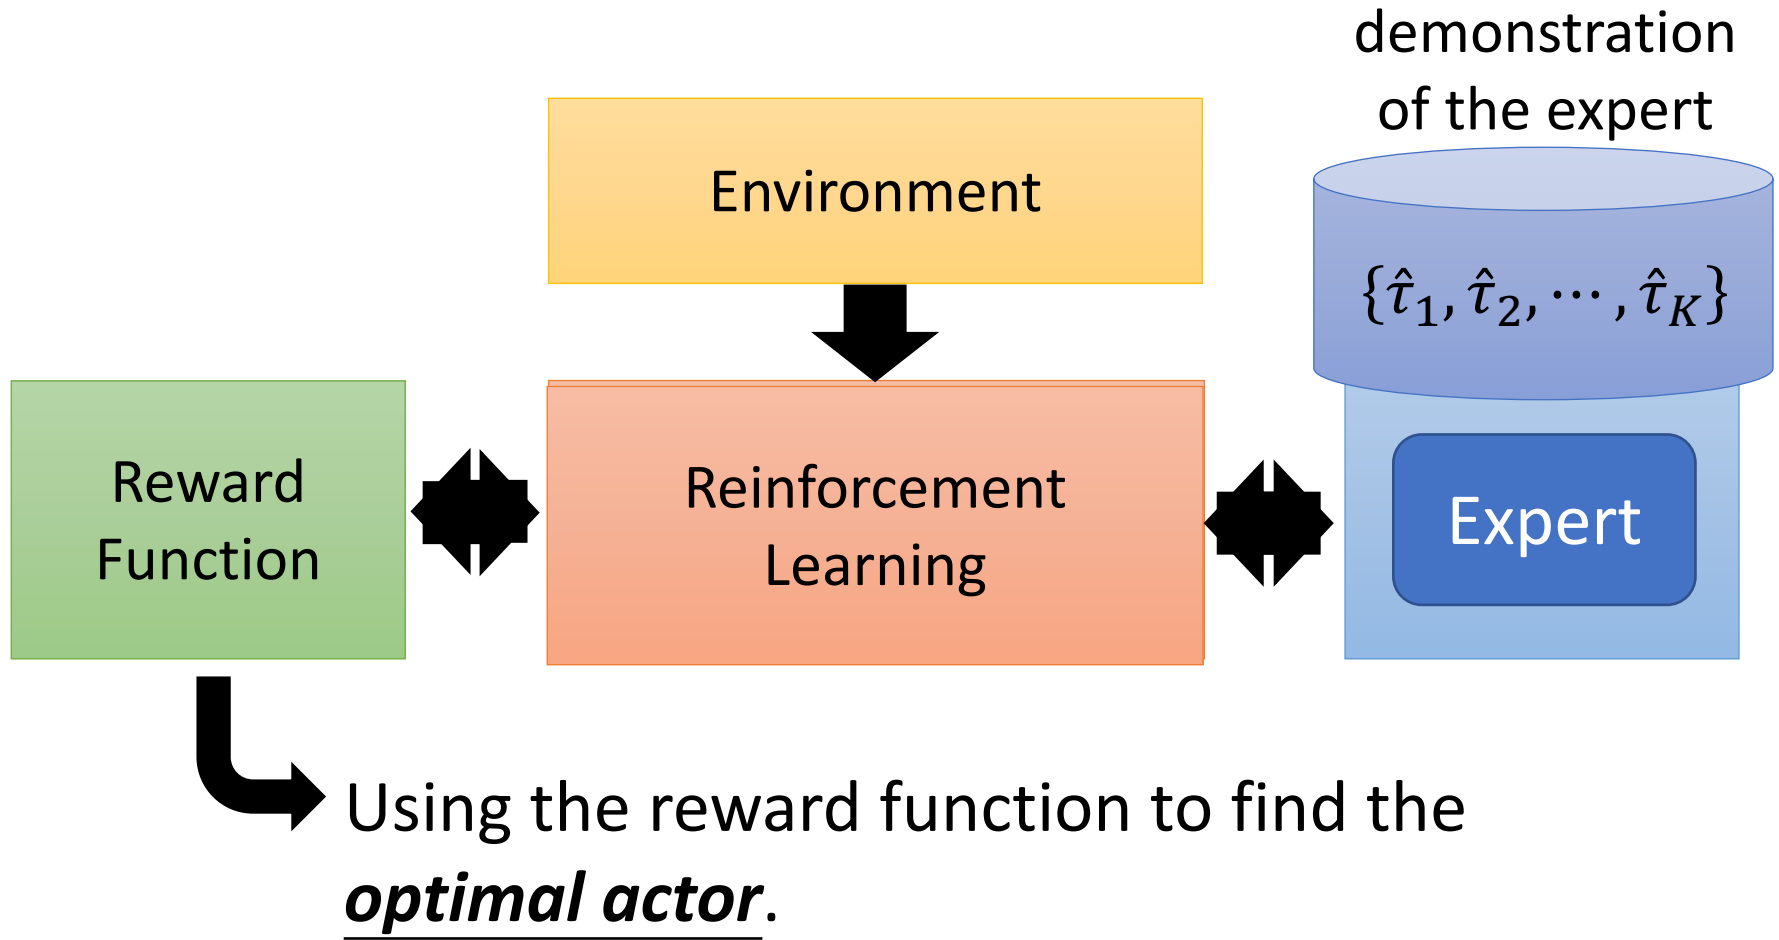
Inverse Reinforcement Learning

- Principle: ***The teacher is always the best***.

- Basic idea:
  - Initialize an actor
  - In each iteration
    - The actor interacts with the environments to obtain some trajectories.
    - Define a reward function, which makes the trajectories of the teacher better than the actor.
    - The actor learns to maximize the reward based on the new reward function.
  - Output the reward function and the actor learned from the reward function
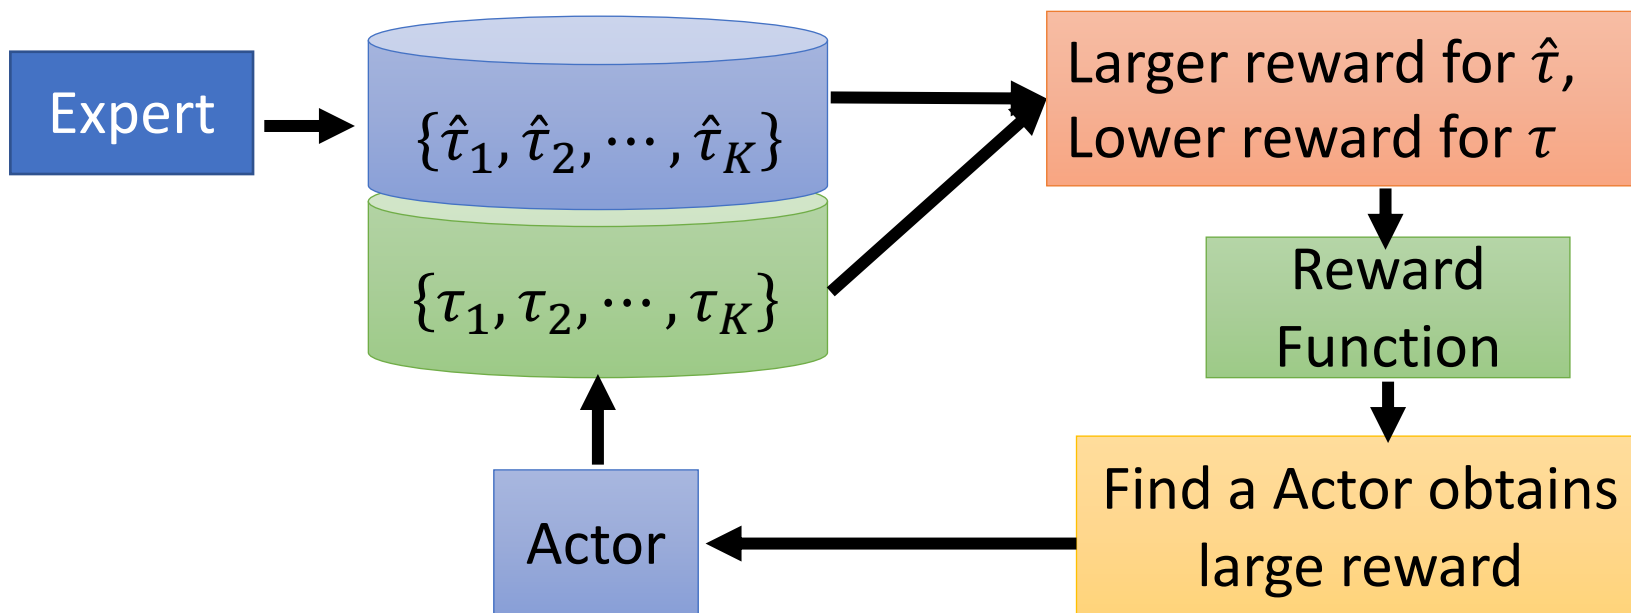
Framework of IRL

$$\sum_{n=1}^{K} R(\hat{\tau}_n) > \sum_{n=1}^{K} R(\tau)$$

Expert $\hat{\pi}$

$\{\hat{\tau}_1, \hat{\tau}_2, \cdots, \hat{\tau}_K\}$

Obtain Reward Function R

$\{\tau_1, \tau_2, \cdots, \tau_K\}$

Reward Function R

Actor
= Generator

Reward function
= Discriminator

Find an actor based on reward function R

Actor π

By Reinforcement learning

GAN

High score for real, low score for generated

D

Find a G whose output obtains large score from D

G

IRL

Expert

$\{\hat{\tau}_1, \hat{\tau}_2, \cdots, \hat{\tau}_K\}$

$\{\tau_1, \tau_2, \cdots, \tau_K\}$

Larger reward for $\hat{\tau}$, Lower reward for τ

Reward Function

Find a Actor obtains large reward

Actor

Robot

- How to teach robots? https://www.youtube.com/watch?v=DEGbtjTOIB0

Robot

Chelsea Finn, Sergey Levine, Pieter Abbeel, Guided Cost Learning: Deep Inverse Optimal Control via Policy Optimization, ICML, 2016
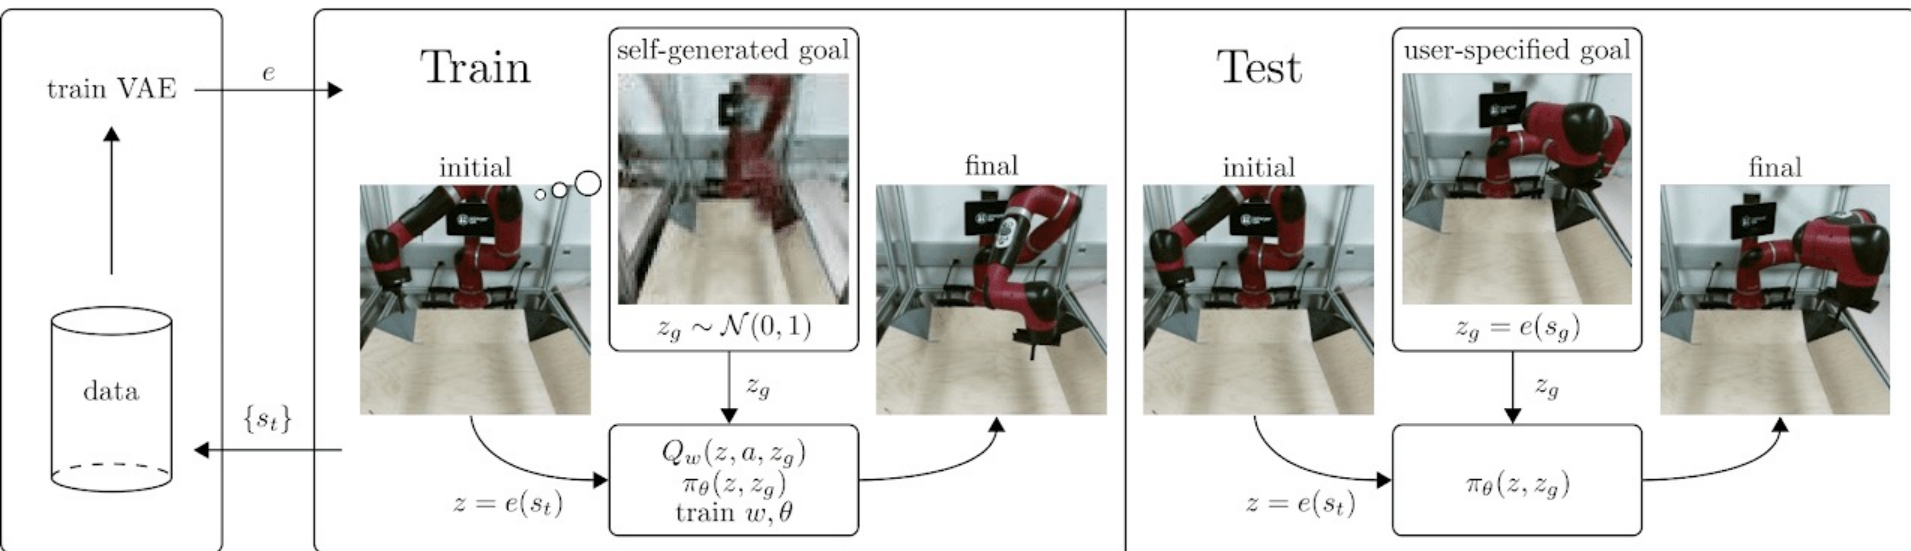
http://rll.berkeley.edu/gcl/

To Learn More …

Visual Reinforcement Learning with Imagined Goals, NIPS 2018

Skew-Fit: State-Covering Self-Supervised Reinforcement Learning, ICML 2020

Reinforcement learning with Imagined Goals (RIG)

Concluding Remarks

What is RL? (Three steps in ML)

Policy Gradient

Actor-Critic

Sparse Reward

No Reward: Learning from Demonstration